

# 知识蒸馏

---

1. KD: Knowledge Distillation
2. FitNet: Hints for thin deep nets
3. AT: Attention Transfer
4. SP: Similarity-Preserving
5. CC: Correlation Congruence
6. VID: Variational Information Distillation
7. RKD: Relation Knowledge Distillation
8. PKT: Probabilistic Knowledge Transfer
9. AB: Activation Boundaries
10. FT: Factor Transfer
11. FSP: Flow of Solution Procedure
12. NST: Neuron Selectivity Transfer
13. CRD: Contrastive Representation Distillation
14. Overhaul

参考文献

[https://blog.csdn.net/qq\\_29462849/article/details/122445107](https://blog.csdn.net/qq_29462849/article/details/122445107)

## 1. KD: Knowledge Distillation

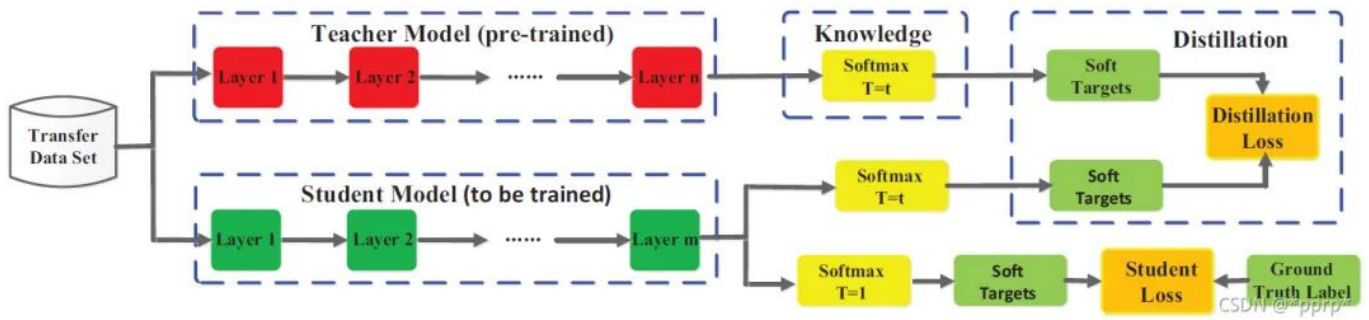
全称: Distilling the Knowledge in a Neural Network

链接: <https://arxiv.org/pdf/1503.02531.pdf>

发表: NIPS14

最经典的，也是明确提出知识蒸馏概念的工作，通过使用带温度的softmax函数来软化教师网络的逻辑层输出作为学生网络的监督信息，

使用KL divergence来衡量学生网络与教师网络的差异，具体流程如下图所示（来自Knowledge Distillation A Survey）



对学生网络来说，一部分监督信息来自hard label标签，另一部分来自教师网络提供的soft label。  
代码实现：

#### 基础KD损失

Plain Text

复制代码

```

1 class DistillKL(nn.Module):
2     """Distilling the Knowledge in a Neural Network"""
3     def __init__(self, T):
4         super(DistillKL, self).__init__()
5         self.T = T
6     def forward(self, y_s, y_t):
7         p_s = F.log_softmax(y_s/self.T, dim=1)
8         p_t = F.softmax(y_t/self.T, dim=1)
9         loss = F.kl_div(p_s, p_t, size_average=False) * (self.T**2) / y_s.shape[0]
10        return loss

```

核心就是一个kl\_div函数，用于计算学生网络和教师网络的分布差异。

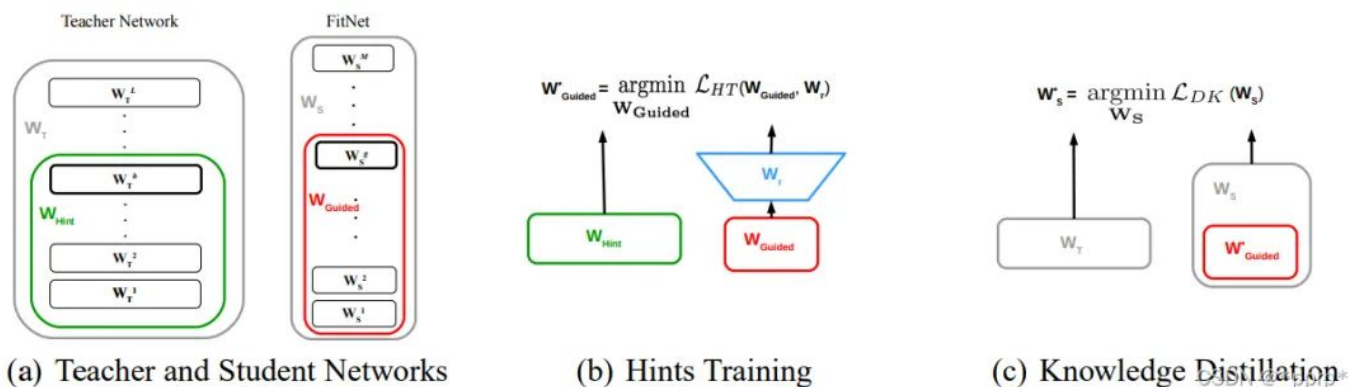
## 2. FitNet: Hints for thin deep nets

全称：Fitnets: hints for thin deep nets

链接：<https://arxiv.org/pdf/1412.6550.pdf>

发表：ICLR 15 Poster

对中间层进行蒸馏的开山之作，通过将学生网络的feature map扩展到与教师网络的feature map相同尺寸以后，使用均方误差MSE Loss来衡量两者差异。



实现如下：

中间特征KD损失

Plain Text

复制代码

```
1 class HintLoss(nn.Module):
2     """Fitnets: hints for thin deep nets, ICLR 2015"""
3     def __init__(self):
4         super(HintLoss, self).__init__()
5         self.crit = nn.MSELoss()
6     def forward(self, f_s, f_t):
7         loss = self.crit(f_s, f_t)
8     return loss
```

实现核心就是MSELoss。


### 3. AT: Attention Transfer

全称：Paying More Attention to Attention: Improving the Performance of Convolutional Neural Networks via Attention Transfer

链接：<https://arxiv.org/pdf/1612.03928.pdf>

发表：ICLR16

为了提升学生模型性能提出使用注意力作为知识载体进行迁移，文中提到了两种注意力，一种是 activation-based attention transfer，另一种是 gradient-based attention transfer。实验发现第一种方法既简单效果又好。



实现如下：

▼ 注意力损失

Plain Text | 复制代码

```
1 class Attention(nn.Module):
2     """Paying More Attention to Attention: Improving the Performance of Convolutional Neural Networks via Attention Transfer code: https://github.com/szagoruyko/attention-transfer"""
3     def __init__(self, p=2):
4         super(Attention, self).__init__()
5         self.p = p
6     def forward(self, g_s, g_t):
7         return [self.at_loss(f_s, f_t) for f_s, f_t in zip(g_s, g_t)]
8     def at_loss(self, f_s, f_t):
9         s_H, t_H = f_s.shape[2], f_t.shape[2]
10        if s_H > t_H:
11            f_s = F.adaptive_avg_pool2d(f_s, (t_H, t_H))
12        elif s_H < t_H:
13            f_t = F.adaptive_avg_pool2d(f_t, (s_H, s_H))
14        else:
15            pass
16        return (self.at(f_s) - self.at(f_t)).pow(2).mean()
17    def at(self, f):
18        return F.normalize(f.pow(self.p).mean(1).view(f.size(0), -1))
```

首先使用avgpool将尺寸调整一致，然后使用MSE Loss来衡量两者差距。


## 4. SP: Similarity-Preserving

全称：Similarity-Preserving Knowledge Distillation

链接：<https://arxiv.org/pdf/1907.09682.pdf>

发表：ICCV19SP

归属于基于关系的知识蒸馏方法。文章思想是提出相似性保留的知识，使得教师网络和学生网络会对相同的样本产生相似的激活。可以从下图看出处理流程，教师网络和学生网络对应feature map通过计算内积，得到bsxbs的相似度矩阵，然后使用均方误差来衡量两个相似度矩阵。



最终Loss为：  
 G代表的就是bsxbs的矩阵。实现如下：

```
1 class Similarity(nn.Module):
2     """Similarity-Preserving Knowledge Distillation, ICCV2019, verified by ori
   ginal author"""
3     def __init__(self):
4         super(Similarity, self).__init__()
5     def forward(self, g_s, g_t):
6         return [self.similarity_loss(f_s, f_t) for f_s, f_t in zip(g_s, g_t)]
7
8     def similarity_loss(self, f_s, f_t):
9         bsz = f_s.shape[0]
10        f_s = f_s.view(bsz, -1)
11        f_t = f_t.view(bsz, -1)
12        G_s = torch.mm(f_s, torch.t(f_s))
13        # G_s = G_s / G_s.norm(2)
14        G_s = torch.nn.functional.normalize(G_s)
15        G_t = torch.mm(f_t, torch.t(f_t))
16        # G_t = G_t / G_t.norm(2)
17        G_t = torch.nn.functional.normalize(G_t)
18        G_diff = G_t - G_s
19        loss = (G_diff * G_diff).view(-1, 1).sum(0) / (bsz * bsz)
20        return loss
```

## 5. CC: Correlation Congruence

全称: Correlation Congruence for Knowledge Distillation

链接: <https://arxiv.org/pdf/1904.01802.pdf>

发表: ICCV19

CC也归属于基于关系的知识蒸馏方法。不应该仅仅引导教师网络和学生网络单个样本向量之间的差异, 还应该学习两个样本之间的相关性, 而这个相关性使用的是Correlation Congruence 教师网络和学生网络相关性之间的欧氏距离。

整体Loss如下:

实现如下:

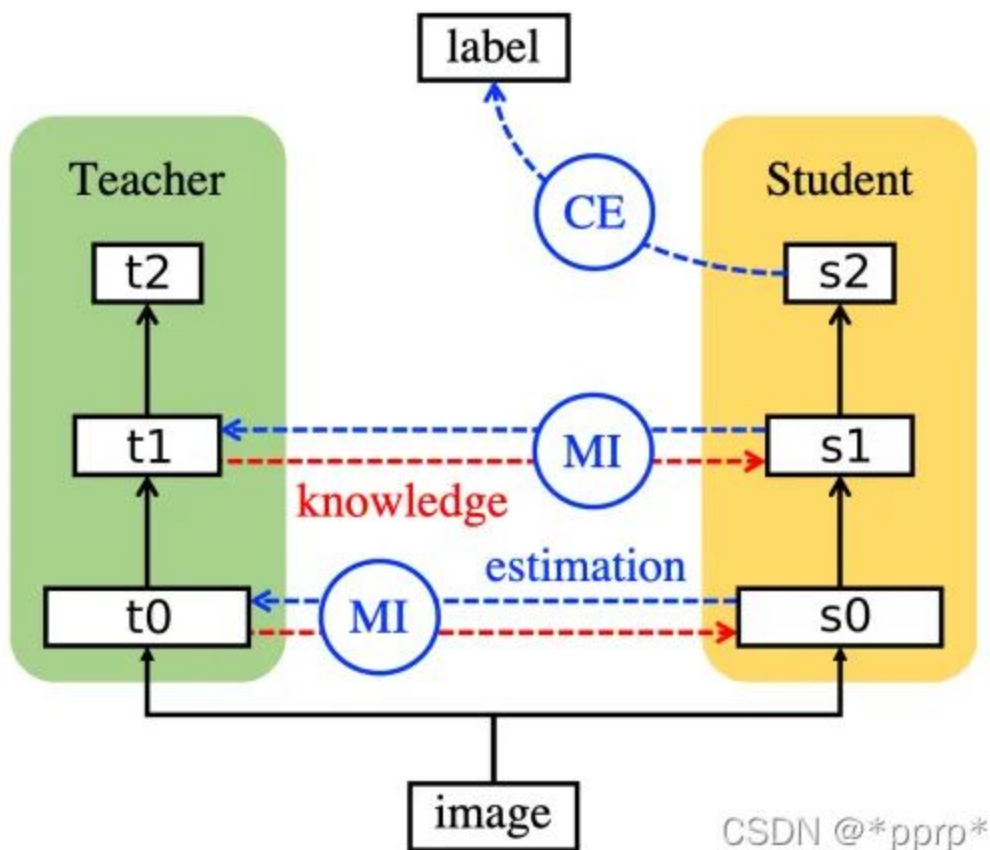
```
1 class Correlation(nn.Module):
2     """Similarity-preserving loss. My original own reimplementation based
3     on the paper before emailing the original authors."""
4     def __init__(self):
5         super(Correlation, self).__init__()
6     def forward(self, f_s, f_t):
7         return self.similarity_loss(f_s, f_t)
8     def similarity_loss(self, f_s, f_t):
9         bsz = f_s.shape[0]
10        f_s = f_s.view(bsz, -1)
11        f_t = f_t.view(bsz, -1)
12        G_s = torch.mm(f_s, torch.t(f_s))
13        G_s = G_s / G_s.norm(2)
14        G_t = torch.mm(f_t, torch.t(f_t))
15        G_t = G_t / G_t.norm(2)
16        G_diff = G_t - G_s
17        loss = (G_diff * G_diff).view(-1, 1).sum(0) / (bsz * bsz)
18        return loss
```

## 6. VID: Variational Information Distillation

全称: Variational Information Distillation for Knowledge Transfer

链接: <https://arxiv.org/pdf/1904.05835.pdf>

发表: CVPR19



利用互信息（Mutual Information）来衡量学生网络和教师网络差异。互信息可以表示出两个变量的互相依赖程度，其值越大，表示变量之间的依赖程度越高。互信息计算如下：

互信息是教师模型的熵减去在已知学生模型条件下教师模型的熵。目标是最大化互信息，因为互信息越大说明 $H(t|s)$ 越小，即学生网络确定的情况下，教师网络的熵会变小，证明学生网络已经学习的比较充分。整体loss如下：

由于 $p(t|s)$ 很难计算，可以使用变分分布 $q(t|s)$ 去接近真实分布。

其中 $q(t|s)$ 是使用方差可学习的高斯分布模拟（公式中的 $\log\_scale$ ）：

实现如下：



```

1 class VIDLoss(nn.Module):
2     """Variational Information Distillation for Knowledge Transfer (CVPR 201
3     9), code from author: https://github.com/ssahn0215/variational-informat
4     ion-distillation"""
5     def __init__(self,
6         num_input_channels,
7         num_mid_channel,
8         num_target_channels,
9         init_pred_var=5.0,
10        eps=1e-5):
11        super(VIDLoss, self).__init__()
12        def conv1x1(in_channels, out_channels, stride=1):
13            return nn.Conv2d(in_channels, out_channels,
14                kernel_size=1, padding=0,
15                bias=False, stride=stride)
16        self.regressor = nn.Sequential(conv1x1(num_input_channels, num_mid_chan
17        nel), nn.ReLU(), conv1x1(num_mid_channel, num_mid_c
18        hannel),
19        nn.ReLU(),
20        conv1x1(num_mid_channel, num_target_channels),)
21        self.log_scale = torch.nn.Parameter(np.log(np.exp(init_pred_var-eps)-
22        1.0) * torch.ones(num_target_channels))
23        self.eps = eps
24        def forward(self, input, target):
25            # pool for dimentsion match
26            s_H, t_H = input.shape[2], target.shape[2]
27            if s_H > t_H:
28                input = F.adaptive_avg_pool2d(input, (t_H, t_H))
29            elif s_H < t_H:
30                target = F.adaptive_avg_pool2d(target, (s_H, s_H))
31            else:
32                pass
33            pred_mean = self.regressor(input)
34            pred_var = torch.log(1.0+torch.exp(self.log_scale))+self.eps
35            pred_var = pred_var.view(1, -1, 1, 1)
36            neg_log_prob = 0.5*((pred_mean-target)**2/pred_var+torch.log(pred_va
37            r))
38            loss = torch.mean(neg_log_prob)
39            return loss

```

## 7. RKD: Relation Knowledge Distillation

全称: Relational Knowledge Disitllation

链接: <http://arxiv.org/pdf/1904.05068>

发表: CVPR19

RKD也是基于关系的知识蒸馏方法, RKD提出了两种损失函数, 二阶的距离损失和三阶的角度损失。

- Distance-wise Loss
- Angle-wise Loss

实现如下:

```

1  class RKDLoss(nn.Module):
2      """Relational Knowledge Disitllation, CVPR2019"""
3      def __init__(self, w_d=25, w_a=50):
4          super(RKDLoss, self).__init__()
5          self.w_d = w_d
6          self.w_a = w_a
7      def forward(self, f_s, f_t):
8          student = f_s.view(f_s.shape[0], -1)
9          teacher = f_t.view(f_t.shape[0], -1)
10         # RKD distance loss
11         with torch.no_grad():
12             t_d = self.pdist(teacher, squared=False)
13             mean_td = t_d[t_d > 0].mean()
14             t_d = t_d / mean_td
15             d = self.pdist(student, squared=False)
16             mean_d = d[d > 0].mean()
17             d = d / mean_d
18         loss_d = F.smooth_l1_loss(d, t_d)
19         # RKD Angle loss
20         with torch.no_grad():
21             td = (teacher.unsqueeze(0) - teacher.unsqueeze(1))
22             norm_td = F.normalize(td, p=2, dim=2)
23             t_angle = torch.bmm(norm_td, norm_td.transpose(1, 2)).view(-1)
24             sd = (student.unsqueeze(0) - student.unsqueeze(1))
25             norm_sd = F.normalize(sd, p=2, dim=2)
26             s_angle = torch.bmm(norm_sd, norm_sd.transpose(1, 2)).view(-1)
27         loss_a = F.smooth_l1_loss(s_angle, t_angle)
28         loss = self.w_d * loss_d + self.w_a * loss_a
29         return loss
30     @staticmethod
31     def pdist(e, squared=False, eps=1e-12):
32         e_square = e.pow(2).sum(dim=1)
33         prod = e @ e.t()
34         res = (e_square.unsqueeze(1) + e_square.unsqueeze(0) - 2 * prod).clamp
(min=eps)
35         if not squared:
36             res = res.sqrt()
37             res = res.clone()
38             res[range(len(e)), range(len(e))] = 0
39         return res

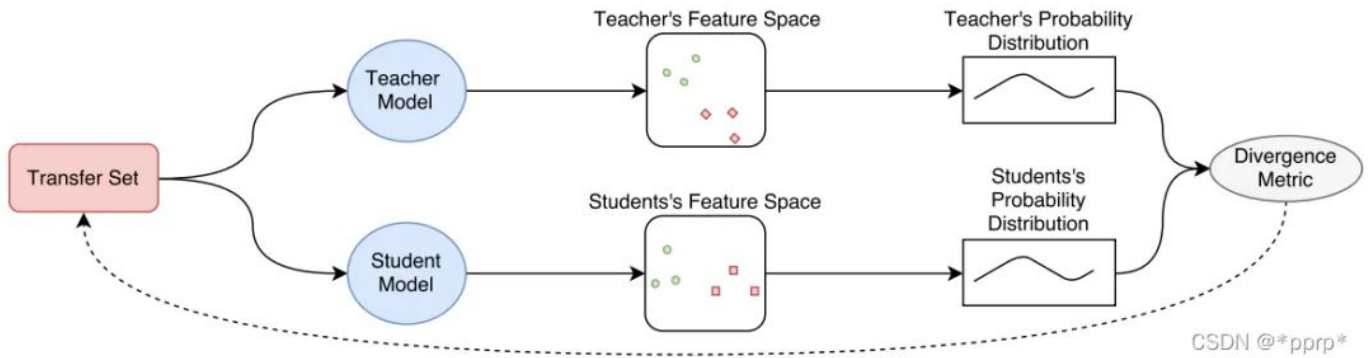
```

## 8. PKT: Probabilistic Knowledge Transfer

全称：Probabilistic Knowledge Transfer for deep representation learning链接：

<https://arxiv.org/abs/1803.10837>发表：CoRR18

提出一种概率知识转移方法，引入了互信息来进行建模。该方法具有可跨模态知识转移、无需考虑任务类型、可将手工特征融入网络等有点。



实现如下：

```
1 class PKT(nn.Module):
2     """Probabilistic Knowledge Transfer for deep representation learning    Co
   de from author: https://github.com/passalis/probabilistic_kt"""
3     def __init__(self):
4         super(PKT, self).__init__()
5     def forward(self, f_s, f_t):
6         return self.cosine_similarity_loss(f_s, f_t)
7     @staticmethod
8     def cosine_similarity_loss(output_net, target_net, eps=0.0000001):
9         # Normalize each vector by its norm
10        output_net_norm = torch.sqrt(torch.sum(output_net ** 2, dim=1, keepdim
   =True))
11        output_net = output_net / (output_net_norm + eps)
12        output_net[output_net != output_net] = 0
13        target_net_norm = torch.sqrt(torch.sum(target_net ** 2, dim=1, keepdim
   =True))
14        target_net = target_net / (target_net_norm + eps)
15        target_net[target_net != target_net] = 0
16        # Calculate the cosine similarity
17        model_similarity = torch.mm(output_net, output_net.transpose(0, 1))
18
19        target_similarity = torch.mm(target_net, target_net.transpose(0, 1))
20        # Scale cosine similarity to 0..1
21        model_similarity = (model_similarity + 1.0) / 2.0
22        target_similarity = (target_similarity + 1.0) / 2.0
23        # Transform them into probabilities
24        model_similarity = model_similarity / torch.sum(model_similarity, dim=
   1, keepdim=True)
25        target_similarity = target_similarity / torch.sum
   (target_similarity, dim=1, keepdim=True)
26        # Calculate the KL-divergence
27        loss = torch.mean(target_similarity * torch.log((target_simila
   rity + eps) / (model_similarity + eps)))
28        return loss
```


## 9. AB: Activation Boundaries

全称: Knowledge Transfer via Distillation of Activation Boundaries Formed by Hidden Neurons

链接: <https://arxiv.org/pdf/1811.03233.pdf>

发表: AAAI18

目标：让教师网络层的神经元的激活边界尽量和学生网络的一样。所谓的激活边界指的是分离超平面（针对的是RELU这种激活函数），其决定了神经元的激活与失活。AB提出的激活转移损失，让教师网络与学生网络之间的分离边界尽可能一致。



实现如下：

#### ▼ 激活损失

Plain Text | 复制代码

```
1 class ABLoss(nn.Module):
2     """Knowledge Transfer via Distillation of Activation Boundaries Formed by
3     Hidden Neurons    code: https://github.com/bhheo/AB_distillation    """
4     def __init__(self, feat_num, margin=1.0):
5         super(ABLoss, self).__init__()
6         self.w = [2*(i-feat_num+1) for i in range(feat_num)]
7         self.margin = margin
8     def forward(self, g_s, g_t):
9         bsz = g_s[0].shape[0]
10        losses = [self.criterion_alternative_l2(s, t) for s, t in zip(g_s, g_t)]
11        losses = [w * l for w, l in zip(self.w, losses)]
12        # loss = sum(losses) / bsz
13        # loss = loss / 1000 * 3
14        losses = [l / bsz for l in losses]
15        losses = [l / 1000 * 3 for l in losses]
16        return losses
17    def criterion_alternative_l2(self, source, target):
18        loss = ((source + self.margin) ** 2 * ((source > -self.margin) & (target <= 0)).float() +
19               (source - self.margin) ** 2 * ((source <= self.margin) & (target > 0)).float())
20        return torch.abs(loss).sum()
```


## 10. FT: Factor Transfer

全称: Paraphrasing Complex Network: Network Compression via Factor Transfer

链接: <https://arxiv.org/pdf/1802.04977.pdf>

发表: NIPS18

提出的是factor transfer的方法。所谓的factor，其实是对模型最后的数据结果进行一个编解码的过程，提取出的一个factor矩阵，用教师网络的factor来指导学生网络的factor。



FT计算公式为：

实现如下：

```
1 class FactorTransfer(nn.Module):
2     """Paraphrasing Complex Network: Network Compression via Factor Transfer,
3     NeurIPS 2018"""
4     def __init__(self, p1=2, p2=1):
5         super(FactorTransfer, self).__init__()
6         self.p1 = p1
7         self.p2 = p2
8     def forward(self, f_s, f_t):
9         return self.factor_loss(f_s, f_t)
10    def factor_loss(self, f_s, f_t):
11        s_H, t_H = f_s.shape[2], f_t.shape[2]
12        if s_H > t_H:
13            f_s = F.adaptive_avg_pool2d(f_s, (t_H, t_H))
14        elif s_H < t_H:
15            f_t = F.adaptive_avg_pool2d(f_t, (s_H, s_H))
16        else:
17            pass
18        if self.p2 == 1:
19            return (self.factor(f_s) - self.factor(f_t)).abs().mean()
20        else:
21            return (self.factor(f_s) - self.factor(f_t)).pow(self.p2).mean()
22    def factor(self, f):
23        return F.normalize(f.pow(self.p1).mean(1).view(f.size(0), -1))
```

## 11. FSP: Flow of Solution Procedure

全称: A Gift from Knowledge Distillation: Fast Optimization, Network Minimization and Transfer Learning

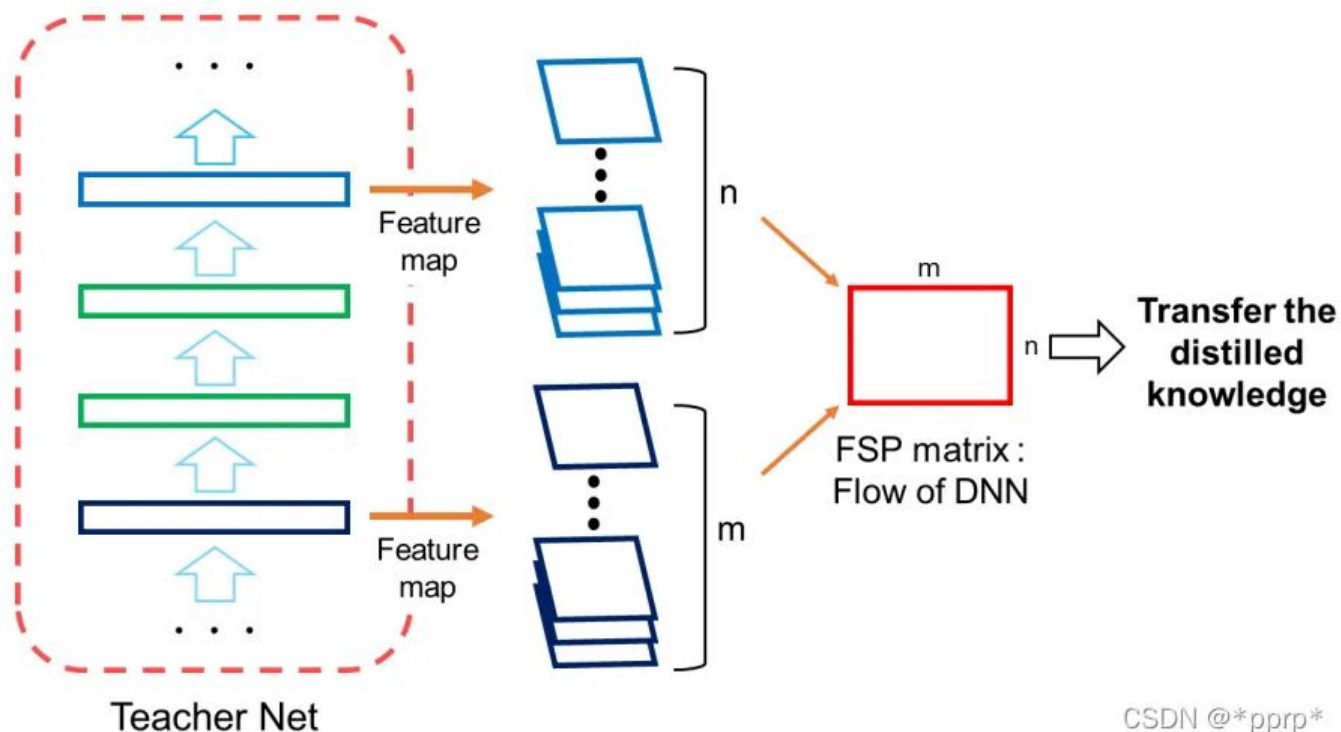
链接:

[https://openaccess.thecvf.com/content\\_cvpr\\_2017/papers/Yim\\_A\\_Gift\\_From\\_CVPR\\_2017\\_paper.pdf](https://openaccess.thecvf.com/content_cvpr_2017/papers/Yim_A_Gift_From_CVPR_2017_paper.pdf)

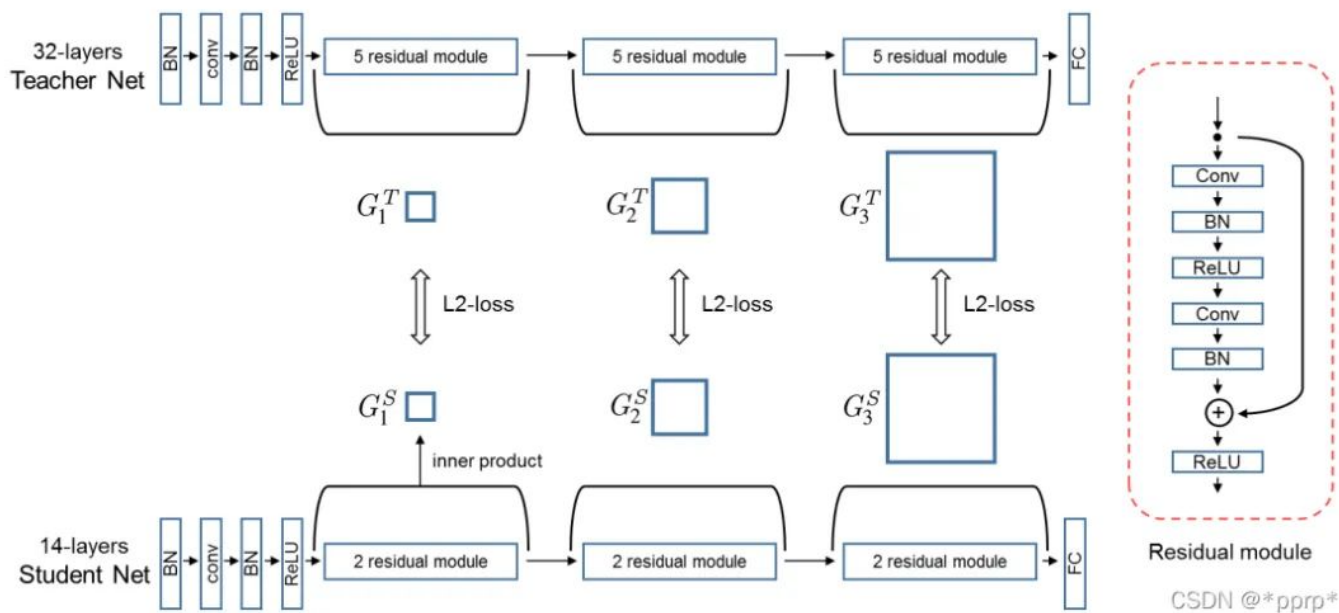
发表: CVPR17

FSP认为教学生网络不同层输出的feature之间的关系比教学生网络结果好





定义了FSP矩阵来定义网络内部特征层之间的关系，是一个Gram矩阵反映老师教学生的过程。



使用的是L2 Loss进行约束FSP矩阵。实现如下：

```
1 class FSP(nn.Module):
2     """A Gift from Knowledge Distillation: Fast Optimization, Network Minimization and Transfer Learning"""
3     def __init__(self, s_shapes, t_shapes):
4         super(FSP, self).__init__()
5         assert len(s_shapes) == len(t_shapes), 'unequal length of feat list'
6
7         s_c = [s[1] for s in s_shapes]
8         t_c = [t[1] for t in t_shapes]
9         if np.any(np.asarray(s_c) != np.asarray(t_c)):
10             raise ValueError('num of channels not equal (error in FSP)')
11
12     def forward(self, g_s, g_t):
13         s_fsp = self.compute_fsp(g_s)
14         t_fsp = self.compute_fsp(g_t)
15         loss_group = [self.compute_loss(s, t) for s, t in zip(s_fsp, t_fsp)]
16
17         return loss_group
18
19     @staticmethod
20     def compute_loss(s, t):
21         return (s - t).pow(2).mean()
22
23     @staticmethod
24     def compute_fsp(g):
25         fsp_list = []
26         for i in range(len(g) - 1):
27             bot, top = g[i], g[i + 1]
28             b_H, t_H = bot.shape[2], top.shape[2]
29             if b_H > t_H:
30                 bot = F.adaptive_avg_pool2d(bot, (t_H, t_H))
31             elif b_H < t_H:
32                 top = F.adaptive_avg_pool2d(top, (b_H, b_H))
33             else:
34                 pass
35             bot = bot.unsqueeze(1)
36             top = top.unsqueeze(2)
37             bot = bot.view(bot.shape[0], bot.shape[1], bot.shape[2], -1)
38             top = top.view(top.shape[0], top.shape[1], top.shape[2], -1)
39             fsp = (bot * top).mean(-1)
40             fsp_list.append(fsp)
41         return fsp_list
```

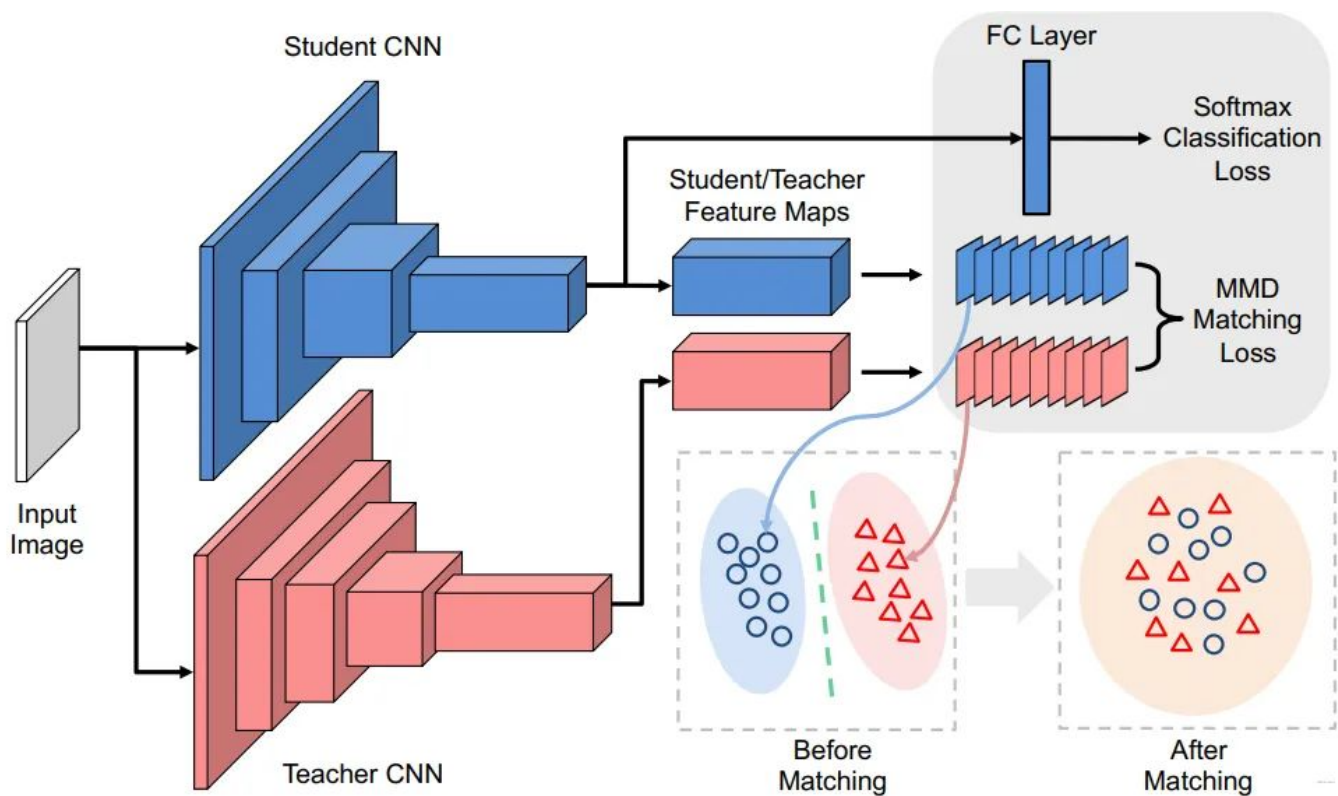
## 12. NST: Neuron Selectivity Transfer

全称: Like what you like: knowledge distill via neuron selectivity transfer

链接: <https://arxiv.org/pdf/1707.01219.pdf>

发表: CoRR17

使用新的损失函数最小化教师网络与学生网络之间的Maximum Mean Discrepancy (MMD), 文中选择的是对其教师网络与学生网络之间神经元选择样式的分布。



使用核技巧(对应下面poly kernel)并进一步展开以后可得:

实际上提供了Linear Kernel、Poly Kernel、Gaussian Kernel三种, 这里实现只给了Poly这种, 这是因为Poly这种方法可以与KD进行互补, 这样整体效果会非常好。实现如下:

```
1 class NSTLoss(nn.Module):
2     """like what you like: knowledge distill via neuron selectivity transfe
3     r"""
4     def __init__(self):
5         super(NSTLoss, self).__init__()
6         pass
7     def forward(self, g_s, g_t):
8         return [self.nst_loss(f_s, f_t) for f_s, f_t in zip(g_s, g_t)]
9     def nst_loss(self, f_s, f_t):
10        s_H, t_H = f_s.shape[2], f_t.shape[2]
11        if s_H > t_H:
12            f_s = F.adaptive_avg_pool2d(f_s, (t_H, t_H))
13        elif s_H < t_H:
14            f_t = F.adaptive_avg_pool2d(f_t, (s_H, s_H))
15        else:
16            pass
17        f_s = f_s.view(f_s.shape[0], f_s.shape[1], -1)
18        f_s = F.normalize(f_s, dim=2)
19        f_t = f_t.view(f_t.shape[0], f_t.shape[1], -1)
20        f_t = F.normalize(f_t, dim=2)
21        # set full_loss as False to avoid unnecessary computation
22        full_loss = True
23        if full_loss:
24            return (self.poly_kernel(f_t, f_t).mean().detach() + self.poly_kerne
25            l(f_s, f_s).mean() - 2 * self.poly_kernel(f_s, f_t).mean())
26        else:
27            return self.poly_kernel(f_s, f_s).mean() - 2 * self.poly_kernel(f_
28            s, f_t).mean()
29        def poly_kernel(self, a, b):
30            a = a.unsqueeze(1)
31            b = b.unsqueeze(2)
32            res = (a * b).sum(-1).pow(2)
33            return res
```

## 13. CRD: Contrastive Representation Distillation

全称: Contrastive Representation Distillation

链接: <https://arxiv.org/abs/1910.10699v2>

发表: ICLR20

将对比学习引入知识蒸馏中，其目标修正为：学习一个表征，让正样本对的教师网络与学生网络尽可能接近，负样本对教师网络与学生网络尽可能远离。构建的对比学习问题表示如下：

整体的蒸馏Loss表示如下：

实现如下：<https://github.com/HobbitLong/RepDistiller>

```

1  class ContrastLoss(nn.Module):
2  """    contrastive loss, corresponding to Eq (18)    """
3      def __init__(self, n_data):
4          super(ContrastLoss, self).__init__()
5          self.n_data = n_data
6      def forward(self, x):
7          bsz = x.shape[0]
8          m = x.size(1) - 1
9          # noise distribution
10         Pn = 1 / float(self.n_data)
11         # loss for positive pair
12         P_pos = x.select(1, 0)
13         log_D1 = torch.div(P_pos, P_pos.add(m * Pn + eps)).log_()
14         # loss for K negative pair
15         P_neg = x.narrow(1, 1, m)
16         log_D0 = torch.div(P_neg.clone().fill_(m * Pn), P_neg.add(m * Pn + eps)).log_()
17         loss = - (log_D1.sum(0) + log_D0.view(-1, 1).sum(0)) / bsz
18         return loss
19  class CRDLoss(nn.Module):
20      """CRD Loss function    includes two symmetric parts:
21          (a) using teacher as anchor, choose positive and negatives over the student side
22          (b) using student as anchor, choose positive and negatives over the teacher side
23          Args:
24          opt.s_dim: the dimension of student's feature
25          opt.t_dim: the dimension of teacher's feature
26          opt.feats_dim: the dimension of the projection space
27          opt.nce_k: number of negatives paired with each positive
28          opt.nce_t: the temperature
29          opt.nce_m: the momentum for updating the memory buffer
30          opt.n_data: the number of samples in the training set, therefore the memory buffer is: opt.n_data x opt.feats_dim """
31      def __init__(self, opt):
32          super(CRDLoss, self).__init__()
33          self.embed_s = Embed(opt.s_dim, opt.feats_dim)
34          self.embed_t = Embed(opt.t_dim, opt.feats_dim)
35          self.contrast = ContrastMemory(opt.feats_dim, opt.n_data, opt.nce_k, opt.nce_t, opt.nce_m)
36          self.criterion_t = ContrastLoss(opt.n_data)
37          self.criterion_s = ContrastLoss(opt.n_data)
38      def forward(self, f_s, f_t, idx, contrast_idx=None):
39          """
40          Args:

```

```

41     f_s: the feature of student network, size [batch_size, s_dim]
42     f_t: the feature of teacher network, size [batch_size, t_dim]
43     idx: the indices of these positive samples in the dataset, size [batch
44         _size]
45     contrast_idx: the indices of negative samples, size [batch_size, nce_
46         k]
47     Returns:
48     The contrastive loss
49     """
50     f_s = self.embed_s(f_s)
51     f_t = self.embed_t(f_t)
52     out_s, out_t = self.contrast(f_s, f_t, idx, contrast_idx)
53     s_loss = self.criterion_s(out_s)
54     t_loss = self.criterion_t(out_t)
55     loss = s_loss + t_loss
56     return loss

```

## 14. Overhaul

全称: A Comprehensive Overhaul of Feature Distillation链接:

[http://openaccess.thecvf.com/content\\_ICCV\\_2019/papers/](http://openaccess.thecvf.com/content_ICCV_2019/papers/)发表: CVPR19

- teacher transform中提出使用margin RELU激活函数。




Figure 4. A comparison of the conventional ReLU, teacher transforms in Heo *et al.* [7] and our proposed method. CSDN @\*pprp\*

- student transform中提出使用1x1卷积。
- distillation feature position选择Pre-ReLU。



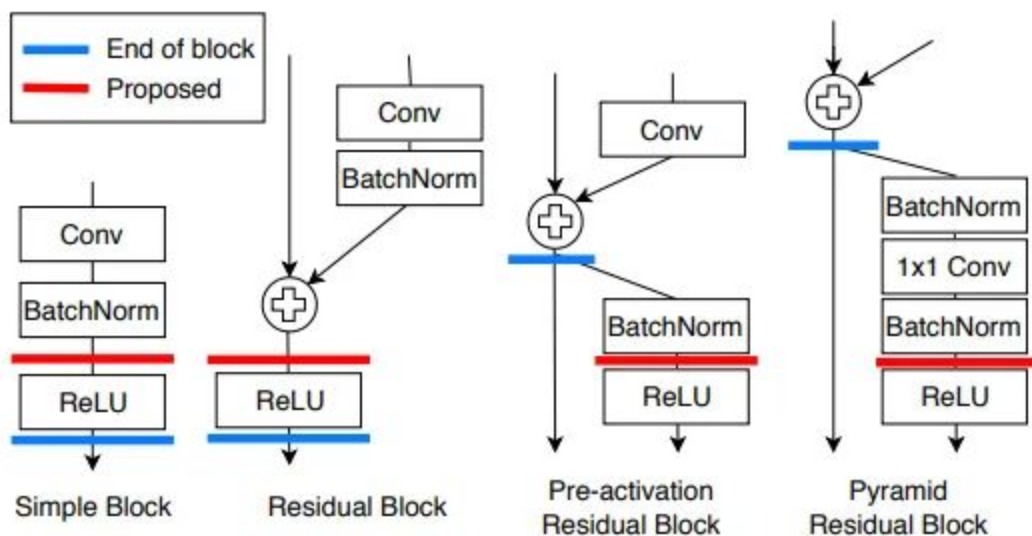


Figure 3. Position of distillation target layer. We place the distillation layer between the last block and the first ReLU. The exact location differs according to the network architecture. CSDN @\*pprp\*

- distance function部分提出了Partial L2 损失函数。

Method	Teacher transform	Student transform	Distillation feature position	Distance	Missing information
FitNets [22]	None	$1 \times 1$ conv	Mid layer	$L_2$	None
AT [30]	Attention	Attention	End of group	$L_2$	Channel dims
FSP [28]	Correlation	Correlation	End of group	$L_2$	Spatial dims
Jacobian [26]	Gradient	Gradient	End of group	$L_2$	Channel dims
FT [13]	Auto-encoder	Auto-encoder	End of group	$L_1$	Auto-encoded
AB [7]	Binarization	$1 \times 1$ conv	Pre-ReLU	Marginal $L_2$	Feature values
Proposed	Margin ReLU	$1 \times 1$ conv	Pre-ReLU	Partial $L_2$	Negative features

Table 1. Difference in various kinds of feature distillation. Most distillation use teacher transform with information loss. CSDN @\*pprp\*

部分实现如下：



```
1 class OFD(nn.Module):
2     '''
3     A Comprehensive Overhaul of Feature Distillation
4     http://openaccess.thecvf.com/content_ICCV_2019/papers/Heo_A_Comprehensiv
5     e_Overhaul_of_Feature_Distillation_ICCV_2019_paper.pdf
6     '''
7     def __init__(self, in_channels, out_channels):
8         super(OFD, self).__init__()
9         self.connector = nn.Sequential(*[
10             nn.Conv2d(in_channels, out_channels, kernel_size=1, stride=1, padding=
11             0, bias=False),
12             nn.BatchNorm2d(out_channels)
13         ])
14         for m in self.modules():
15             if isinstance(m, nn.Conv2d):
16                 nn.init.kaiming_normal_(m.weight, mode='fan_out', nonlinearity='re
17                 lu')
18             if m.bias is not None:
19                 nn.init.constant_(m.bias, 0)
20             elif isinstance(m, nn.BatchNorm2d):
21                 nn.init.constant_(m.weight, 1)
22                 nn.init.constant_(m.bias, 0)
23         def forward(self, fm_s, fm_t):
24             margin = self.get_margin(fm_t)
25             fm_t = torch.max(fm_t, margin)
26             fm_s = self.connector(fm_s)
27             mask = 1.0 - ((fm_s <= fm_t) & (fm_t <= 0.0)).float()
28             loss = torch.mean((fm_s - fm_t)**2 * mask)
29             return loss
30         def get_margin(self, fm, eps=1e-6):
31             mask = (fm < 0.0).float()
32             masked_fm = fm * mask
33             margin = masked_fm.sum(dim=(0,2,3), keepdim=True) / (mask.sum(dim=(0,
34             2,3), keepdim=True)+eps)
35             return margin
```

## 参考文献

[https://blog.csdn.net/weixin\\_44579633/article/details/119350631](https://blog.csdn.net/weixin_44579633/article/details/119350631)

<https://blog.csdn.net/winycg/article/details/105297089>

[https://blog.csdn.net/weixin\\_46239293/article/details/120289163](https://blog.csdn.net/weixin_46239293/article/details/120289163)

[https://blog.csdn.net/DD\\_PP\\_JJ/article/details/121578722](https://blog.csdn.net/DD_PP_JJ/article/details/121578722)

[https://blog.csdn.net/DD\\_PP\\_JJ/article/details/121714957](https://blog.csdn.net/DD_PP_JJ/article/details/121714957)

<https://zhuanlan.zhihu.com/p/344881975>

[https://blog.csdn.net/weixin\\_44633882/article/details/108927033](https://blog.csdn.net/weixin_44633882/article/details/108927033)

[https://blog.csdn.net/weixin\\_46239293/article/details/120266111](https://blog.csdn.net/weixin_46239293/article/details/120266111)

[https://blog.csdn.net/weixin\\_43402775/article/details/109011296](https://blog.csdn.net/weixin_43402775/article/details/109011296)

[https://blog.csdn.net/m0\\_37665984/article/details/103288582](https://blog.csdn.net/m0_37665984/article/details/103288582)

[https://blog.csdn.net/m0\\_37665984/article/details/103269740](https://blog.csdn.net/m0_37665984/article/details/103269740)